

Model Optimization and Tuning Phase Report

Date	: 22-06-2025
Team ID	: SWTID1749705847
Project Title	: “ Rising Waters: A Machine Learning Approach to Flood Prediction”
Maximum Marks	: 10 Marks

Hyper-parameter Tuning

Documentation:

MODEL	Tuned Hyperparameters
Decision Tree	<pre># Decision Tree with GridSearchCV from sklearn.tree import DecisionTreeClassifier dt_classifier = DecisionTreeClassifier() param_grid_dt = { 'criterion': ['gini', 'entropy'], 'splitter': ['best', 'random'], 'max_depth': [None, 10, 20, 30, 40, 50], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4] }</pre>
Random Forest	<pre># Random Forest with GridSearchCV from sklearn.ensemble import RandomForestClassifier rf_classifier = RandomForestClassifier() param_grid_rf = { 'n_estimators': [50, 100, 200], 'criterion': ['gini', 'entropy'], 'max_depth': [None, 10, 20, 30], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4] }</pre>

KNN	<pre># KNN with GridSearchCV from sklearn.neighbors import KNeighborsClassifier knn_classifier = KNeighborsClassifier() param_grid_knn = { 'n_neighbors': [3, 5, 7, 9], 'weights': ['uniform', 'distance'], 'p': [1, 2] }</pre>
Gradient Boosting	<pre># Gradient Boosting with GridSearchCV from sklearn.ensemble import GradientBoostingClassifier gb_classifier = GradientBoostingClassifier() param_grid_gb = { 'n_estimators': [50, 100, 200], 'learning_rate': [0.01, 0.1, 0.2], 'max_depth': [3, 4, 5], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4], 'subsample': [0.8, 1.0] }</pre>

Optimal Values

```
grid_dt = GridSearchCV(estimator=dt_classifier, param_grid=param_grid_dt, cv=5)
grid_dt.fit(x_train, y_train)
y_pred_dt = grid_dt.predict(x_test)
print("Optimal Hyperparameters:", grid_dt.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred_dt))
```

Optimal Hyperparameters: {'criterion': 'entropy', 'max_depth': 30, 'min_samples_leaf': 2, 'min_samples_split': 2, 'splitter': 'random'}
Accuracy: 1.0

```
grid_rf = GridSearchCV(estimator=rf_classifier, param_grid=param_grid_rf, cv=5)
grid_rf.fit(x_train, np.ravel(y_train))
y_pred_rf = grid_rf.predict(x_test)
print("Optimal Hyperparameters:", grid_rf.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
```

Optimal Hyperparameters: {'criterion': 'gini', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}
Accuracy: 0.9615384615384616

```

grid_rf = GridSearchCV(estimator=rf_classifier, param_grid=param_grid_rf, cv=5)
grid_rf.fit(x_train, np.ravel(y_train))
y_pred_rf = grid_rf.predict(x_test)
print("Optimal Hyperparameters:", grid_rf.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred_rf))

```

Optimal Hyperparameters: {'n_neighbors': 5, 'p': 1, 'weights': 'uniform'}
Accuracy: 0.9615384615384616

```

grid_knn = GridSearchCV(estimator=knn_classifier, param_grid=param_grid_knn, cv=5)
grid_knn.fit(x_train, np.ravel(y_train))
y_pred_knn = grid_knn.predict(x_test)
print("Optimal Hyperparameters:", grid_knn.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred_knn))

```

Optimal Hyperparameters: {'learning_rate': 0.2, 'max_depth': 3, 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 50, 'subsample': 0.5}
Accuracy: 1.0

Performance Metrics Comparision Report:

Model	Optimized Metric
Decision Tree	<pre> print(classification_report(y_test, y_pred_dt)) print(confusion_matrix(y_test, y_pred_dt)) </pre> <pre> precision recall f1-score support 0 1.00 1.00 1.00 24 1 1.00 1.00 1.00 2 accuracy 1.00 26 macro avg 1.00 26 weighted avg 1.00 26 [[24 0] [0 2]] </pre>

Random Forest

```
print(classification_report(y_test, y_pred_rf))
print(confusion_matrix(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	24
1	1.00	0.50	0.67	2
accuracy			0.96	26
macro avg	0.98	0.75	0.82	26
weighted avg	0.96	0.96	0.96	26


```
[[24  0]
 [ 1  1]]
```

KNN

```
print(classification_report(y_test, y_pred_knn))
print(confusion_matrix(y_test, y_pred_knn))
```

	precision	recall	f1-score	support
0	1.00	0.96	0.98	24
1	0.67	1.00	0.80	2
accuracy			0.96	26
macro avg	0.83	0.98	0.89	26
weighted avg	0.97	0.96	0.96	26


```
[[23  1]
 [ 0  2]]
```

Gradient Boosting

```
print(classification_report(y_test, y_pred_gb))
print(confusion_matrix(y_test, y_pred_gb))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	24
1	1.00	1.00	1.00	2
accuracy			1.00	26
macro avg	1.00	1.00	1.00	26
weighted avg	1.00	1.00	1.00	26
[[24 0] [0 2]]				

Final Model Selection Justification:

Our Dataset is imbalanced.

Model	Accuracy	Recall (1)	F1 (1)
Decision Tree	1.00	1.00	1.00
Random Forest	0.96	0.50	0.67
KNN	0.96	1.00	0.80
Gradient Boosting	1.00	1.00	1.00

After tuning and evaluating all models, **Gradient Boosting Classifier** was selected as the final model.

- It achieved **100% accuracy** and **perfect scores (1.00)** in all metrics (precision, recall, F1) for both flood and non-flood classes.
- It **handled class imbalance** well and ensured **zero false negatives**, which is crucial in flood detection.
- Compared to Random Forest and KNN, it performed better on the rare flood cases.
- Unlike Decision Tree, it is **less prone to overfitting** and combines multiple learners for stronger performance.
- Gradient Boosting provides **better control over learning rate and depth**, making it ideal for **future improvements**.

In real-world flood prediction, **missing a flood is riskier than a false alarm**, and this model gives the **most reliable and safe predictions**.

